

Rubric Self-Check

Mapping this submission against the capstone brief's stated requirements

This maps every explicit requirement in the project brief to where it's satisfied in this submission. Status is self-assessed honestly: **Done** means real, verifiable evidence exists; **Partial** means work exists but has a named gap; nothing is marked Done without a specific file or test result behind it.

Required Repository Contents

Requirement	Evidence	Status
README with 15-min fork-and-run path	README.md, full setup + quick-start steps	Done
Containerized agent harness	docker/Dockerfile, docker-compose.yml, non-root user, scoped mounts	Done
Module 1: PRD, rubric, baseline, iteration log	docs/prd.md, quality-rubric.md, baseline-metrics.md (real results), iteration-log.md	Done
Module 2: agent def, 2+ skills, memory layout, reflection log	agents/*.md (3), skills/*.md (2) + risk_scoring.py, memory/architecture-notes.md, reflection-log.md	Done
Module 3: orchestration diagram, tool-grant map, MCP config, schemas, retrieval report, eval harness, holdout set, calibration log	docs/orchestration-diagram.md, .mcp.json, mcp/*.py (schemas inline), docs/retrieval-quality-report.md (real 4/6 result), evals/*, memory/store/calibration-log.jsonl	Done
Module 4: governance policy, role-to-tool matrix, CI/CD guardrails, audit log template, failure-mode playbook, decision matrix, deterministic conversion, ADR	docs/governance-policy.md, .github/workflows/policy-checks.yml, docs/audit-log-template.md, docs/decision-matrix.md, skills/risk_scoring.py, docs/adr/ADR-001	Done
Production-like integration: agentic+deterministic+human, monitoring, reliability, cost controls, incident runbook	orchestrator.py, docs/reliability-cost-controls.md (real tests), docs/ops-runbook.md/pdf	Done
Sanitization: no secrets/PII/proprietary data	No real secrets in repo; docker/.env gitignored; all data from public repos	Done

Delivery Path Requirements (Job-Seeker / No-Deployment)

Requirement	Evidence	Status
Run on public/representative dataset	chalk/chalk + date-fns/date-fns real merged PRs, pulled via fetch-prs.sh	Done
Document constraints honestly	Every doc's "Honest Limitations" / "Known Limitations" section (impact-study.md, retrieval-quality-report.md, etc.)	Done
Explain what would change in real deployment	README "Delivery Path" section, ops-runbook "Next Steps for Real Deployment"	Done
Tool-evolution drill	docs/tool-evolution-drill.md — real permission revocation, caught, confirmed, rolled back	Done
Impact report	docs/impact-study.md + impact-report.pdf, real baseline vs. holdout results	Done
ADRs for durable decisions	docs/adr/ ADR-001 through ADR-006	Done
Demonstrate governance stopping an over-reaching agent	tool-evolution-drill.md's live PermissionError; walkthrough video will show this live	Partial

Evaluation and Impact Requirements

Requirement	Evidence	Status
Baseline comparison	docs/impact-study.md Section 3, side-by-side table	Done
5 required metrics (quality, latency, defect rate, cycle time, cost), proxy explained where needed	impact-study.md — quality/defect rate real; latency/cost explicitly flagged as not-yet-comparable pending a fully automated run	Partial
Evaluation harness: deterministic + rubric scoring, holdout results, calibration log	evals/run_eval.py (tested), evals/holdout-set.md, memory/store/calibration-log.jsonl. Full 0-2 rubric scoring of real outputs not yet applied.	Partial
Tool-evolution drill: what broke, what improved, what harness caught, fix, result	docs/tool-evolution-drill.md — all 5 elements present with real captured output	Done
Run evidence: logs, transcripts, commit history, eval outputs	evals/sample-runs/ (real reviewer output), evals/eval-results.jsonl, test_policy.py output	Done

Architecture and Decision Requirements

Requirement	Evidence	Status
Architecture write-up covering workflow, roles, deterministic steps, human checkpoints, data flow, MCP, storage, retrieval, CI/CD, governance	architecture-write-up.pdf, all 11 sections present	Done
ADRs: rubric, memory, MCP boundaries, subagent scoping, governance, 1+ deterministic conversion	docs/adr/ADR-001 through ADR-006 (6 total, all required topics covered)	Done
ADR quality bar: context/decision/rejected alternatives/evidence/consequences/open risks; 1+ rejected alternative cites concrete evidence	All 6 ADRs follow this structure; every ADR has at least one evidence-cited rejected alternative	Done
Right-tool rationale for every major step	docs/decision-matrix.md — all 7 pipeline steps classified with reasoning	Done

Governance and Safety Requirements

Requirement	Evidence	Status
Least privilege per role	docs/governance-policy.md matrix; every agent def's own permission table	Done
Enforcement in code, not just policy doc	mcp/*.py TOOL_GRANTS + tests/test_policy.py (5/5 real passing tests, imports actual server code)	Done
Auditability: log actions, tool calls, denials, approvals, failures, retries, rollbacks	docs/audit-log-template.md (schema + example); orchestrator.py's audit_events list captures this at runtime	Done
Escalation and rollback rules defined	docs/governance-policy.md, ops-runbook.md Sections 3-4 (real, tested rollback procedure)	Done
Security hygiene: secrets handling, data classification, retrieval citations, schema validation, policy-bypass checks	docker/.env gitignored; classification table in governance-policy.md; retrieval results include citation field; storage_server.py schema validation; test_policy.py denial tests	Done

Presentation Requirements

Requirement	Evidence	Status
5-10 min technical walkthrough video	Not yet recorded	Pending
Video shows: workflow/value, repo structure, containerized run, delegation + tool grants + retrieval + eval, governance stopping overreach, decision matrix, impact metrics, tool-evolution summary	Content exists in written form across all deliverables above; video assembly pending	Pending
Submission links + files in Canvas	This document set is the file package; links pending final submission step	Pending

Summary: All required repository contents, delivery-path requirements, evaluation/impact requirements, architecture/decision requirements, and governance/safety requirements have real, file-backed evidence. Two categories remain honestly marked Partial (full 5-metric baseline comparability, pending a fully automated orchestrator run with a live API key; full rubric scoring of real agent output) and one category is Pending (the presentation/video deliverables, not yet recorded).